

PSY 653 Module 07: ARIMA Modeling

Description of the dataset:

The covid2.csv is a time series dataset providing the number of new cases of COVID within the United States. There are three variables:

- **location** = location of event (filtered to only United States)
- **date** = date of the new COVID cases
- **new_cases** = number of new COVID cases on the corresponding day

1. Download the “covid2.csv” dataset from the module 07 lab module on canvas
2. Create a new R notebook from your project file and name it “ARIMA modeling notebook”
3. Create a first level header: “Load Libraries”
 - a. In a new R chunk load in the tseries, forecast, psych & tidyverse packages
4. Create an R-chunk with the first level header: “Import data”
 - a. Read in the datafile “covid2.csv” save it to an object named “covid”
 - i. Note: make sure to use read_csv() (instead of read.csv()) so the date variable will be read as a date (You will *not* need to indicate the date format [i.e. %d/%m/%Y] like we did in lab, read_csv() should read it in automatically).
5. Create a new first level header: “Get variable descriptives”
 - a. Use any method to get the dataset descriptives
6. Create a first level header: “Plot the data”
 - a. Plot a line graph so that the date is on the x-axis and new_cases is on the y-axis
 - i. Hint: ggplot(data, aes(x = x-variable, y = y-variable)) + geom_line()
 - b. Does the data look stationary? Why do you think that?
7. Convert the new_cases variable to a timeseries object using ts(), and add a frequency of 30. Name the new time series object as “count_cases”
 - a. Hint: object <- ts(data\$var, frequency = 30)
8. Use adf.test() to statistically test if your time series object is stationary.
 - a. Hint: adf.test(object)
 - b. Statistically speaking, is your data stationary?
9. Smooth your data using the stl() function with an s.window argument set to “periodic”. Name the new object “decomp”.
 - a. Hint: decomp <- stl(count_cases, s.window = “periodic”)

10. Deseasonalize your data using the `seasadj()` function. Name the new object “deseasonal_cnt”
 - a. Hint: `deseasonal_cnt <- seasadj(decomp)`
11. Take the deseasonalized data and give it a difference of 1. Name the new object “count_d1”
 - a. Hint: `count_d1 <- diff(deseasonal_cnt, difference = 1)`
12. Plot the newly seasonally adjusted data using the `plot()` function
 - a. Hint: `plot(count_d1)`
 - b. How does it look now? Does the differenced data look stationary?
13. Re-check the stationarity of the data using the `adf.test()` function
 - a. Hint: `adf.test(count_d1)`
 - b. Is the data now stationary? How do you know?
14. Create an autocorrelation plot of your `count_d1` object using the `Acf()` function.
 - a. Hint: `Acf(count_d1)`
 - b. Looking at this plot: How many Autocorrelations terms seem like they should be included in the model?
15. Create a partial autocorrelation plot of your `count_d1` object using the `Pacf()` function.
 - a. Hint: `Pacf(count_d1)`
 - b. Looking at this plot: How many Moving average terms seem like they should be included in the model?
16. Use the `arima()` function and fit, what you view is, the best fitting arima model.
 - a. This may take a few steps and different people may come up with different models. That is okay!
17. Use the `auto.arima()` function to have R fit the best arima model.
 - a. Did you and the `auto.arima()` function come to the same conclusions?
18. Plot a forecast of the model!
 - a. Hint: `autoplot(forecast(deseasonal_cnt))`
19. Take a step back and view the predicted forecast you just created. In layman’s terms, What did we just create with this model? What is ARIMA modeling good for?